

11/5/25: Evaluating Launch Darkly AI Configs

Glossary

- Creating AI Model Configs
- Creating AI Configs
- Get SDK Key
- Tracking Metrics
- Experiments
- Evaluations
- Observability
- Tracing
- Creating Tools
- TODO
- POC Code
- Analysis

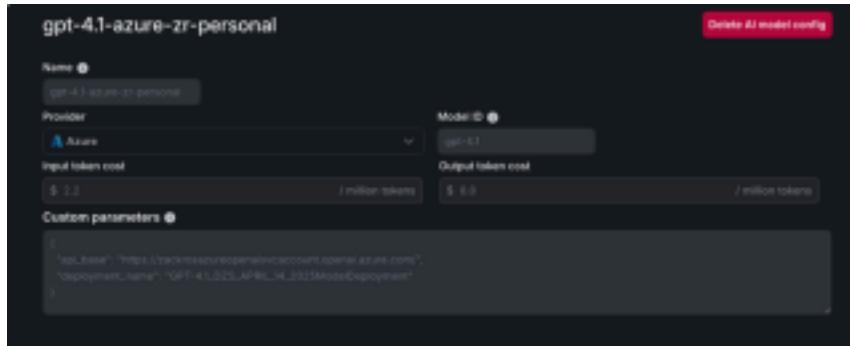
AI Model Config

Config for a particular model deployment

endpoint

deployment name

temperature (maybe – we have the option to override the temperature depending on the use case)

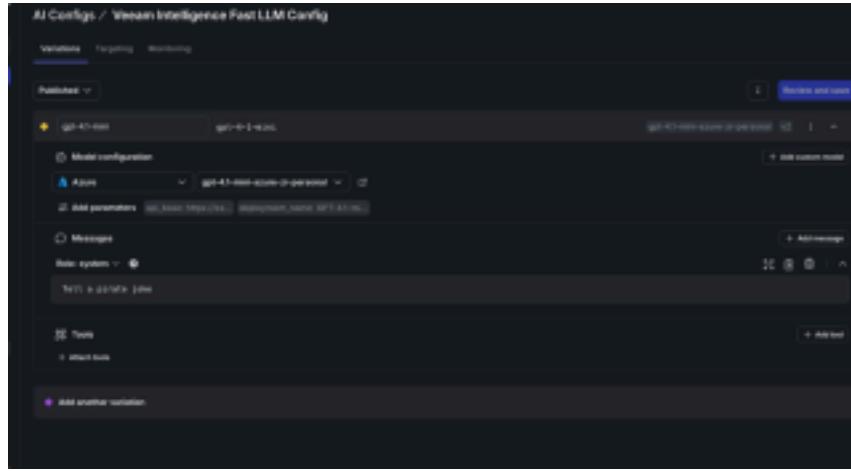


AI Config

Config for a particular LLM-based query

Which AI model config to use

Prompts to be fed to the AI model config



Creating AI Model Configs

First, let's create AI model configs that correspond to the model deployments in our Azure Open AI Service account.

Example model deployment details from Azure Open AI Service

The screenshot displays two separate pages from the Azure Open AI Service portal under the 'Deployments' section:

- SP1-4-1205 APRIL 14, 2023 (Read-Only Deployment)**: This page shows a deployment named 'gpt-4-1205' with a status of 'Read-Only'. It includes sections for 'Get Started' (with steps for authentication and API key), 'Install dependencies' (using pip install openai), and 'Run a basic code sample' (with a provided Python script).
- main endpoint (3 small AI ModelDeployment)**: This page lists three model deployments: 'gpt-4-1205', 'gpt-4-1205-01', and 'gpt-4-1205-02'. Each deployment has its own 'Get Started' section with similar instructions for authentication and running code samples.

In this, it seems AWS Bedrock has a standard endpoint per region and the region needs to be pulled in from the [example code](#) env vars_(it's not part of the Launch Darkly Provider class. Azure Open AI Service does things different than AWS [here](#)

Bedrock – each Azure Open AI Service account has its own endpoint.

We have model deployments in each VDC environment
(personal, dev, stage, prod). See

[Azure OpenAI Configuration | Quot](#)

[a Allocations for Each Environment](#)

for more details. We'll need to create AI model configs for the model deployments in each

env because Azure Open AI Service creates a unique endpoint for each Azure Open AI Service account. We'll add the api base endpoint, the deployment name, and maybe even the api version as a Custom Parameter.

The `endpoint` and `deployment_name` params would probably be a better fit in Model Params rather than Custom Params because these are core params for all of our Azure LLM deployments.

After talking to Launch Darkly solutions engineers, they said this was because AWS Bedrock has a consistent endpoint for all models deployed in the same region (see this).

[example code](#)

Not ideal but at the same time, it's hard to generalize across all model providers and the Custom Params option gives us the flexibility to add the info we need for each AI Model Config.

The image contains three vertically stacked screenshots of the AI Model Config interface, likely from a web application or API configuration tool.

- Screenshot 1: Add custom model**

This is a modal window titled "Add custom model". It includes fields for "Name" (with a placeholder), "Provider" (set to "Azure"), "Model ID" (placeholder), "Input token cost" (\$0.00), and "Output token cost" (\$0.00). A note below says: "Input your parameters for the model. These parameters are available when retrieving a model through the SDK." There are two sections for "Model parameters" and "Custom parameters", both with "+ Add suggested parameters" and "+ Add custom parameters" buttons. A blue "Save" button is at the bottom.
- Screenshot 2: gpt-4.1-azure-zr-personal**

This is a detailed view of a model config named "gpt-4.1-azure-zr-personal". It shows the same basic information as the first screenshot, plus a "Custom parameters" section containing the JSON object: `{"api_base": "https://chattozureopenaiaccount.openai.azure.com", "deployment_name": "GPT-4.1-025-API-001-14_2025ModelDeployment"}`.
- Screenshot 3: gpt-4.1-mini-azure-zr-personal**

This is another detailed view of a model config named "gpt-4.1-mini-azure-zr-personal". It shows the same basic information, plus a "Custom parameters" section containing the JSON object: `{"api_base": "https://chattozureopenaiaccount.openai.azure.com", "deployment_name": "GPT-4.1-mini-025-API-001-14_2025ModelDeployment"}`.

Some AI model providers support “model params” and “custom params” (e.g. Azure) while others only support “custom params” (e.g. Open AI)?

Add custom model

Name:

Provider: OpenAI

Model ID:

Input token cost: \$ 0 / million tokens

Output token cost: \$ 0 / million tokens

Custom parameters: + Add custom parameters

Save

(minor feedback) why can't we edit the Model Config after creation? Have to delete and recreate if we want to change the name, token input/output cost, or the custom params

gpt-4.1

Name: gpt-4.1

Provider: OpenAI

Model ID: gpt-4.1

Input token cost: \$ 0 / million tokens

Output token cost: \$ 0 / million tokens

Custom parameters:

Delete

Creating AI Configs

I created two “completion based” AI config (rather than “AI agent” AI config) because this is the best fit for the Veeam Intelligence workflow.

[Details on “Completion-Based” AI configs vs. “Agent-Based” AI configs](#)

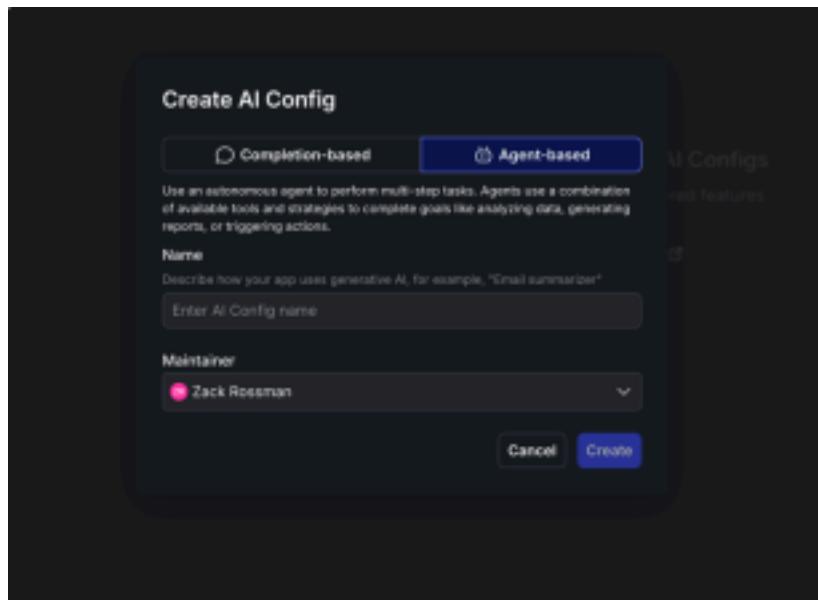
[Quickstart for AI Configs | LaunchDarkly | Documentation](#)
AI Configs can be created in two modes:

Prompt-based: Configure prompts with messages and roles.

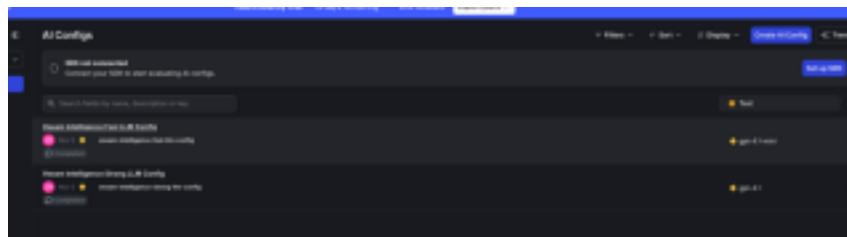
Agent-based: Configure multi-step workflows with instructions and tools. To learn more, read Agents in AI Configs.

[Agents in AI Configs | LaunchDarkly | Documentation](#)

Use agent-based AI Configs to define, monitor, and adapt structured model workflows in real time. Agents use the same configuration model as standard AI Configs, including environment targeting, variation management, and monitoring. The difference is that agents define how a model performs reasoning and takes actions across multiple steps instead of responding to a single prompt.



One AI config for the “strong” model for handling more complex queries, and another for the “fast” model for handling simpler queries.



Get SDK Key

This was tricky. At first I thought the application code needed to use an API key to fetch the AI configs. But actually we need to use an SDK key to fetch the AI configs.

Tracking Metrics

I added a wrapper to Veeam Intelligence to track various metrics

Generation count - Number of successful completions

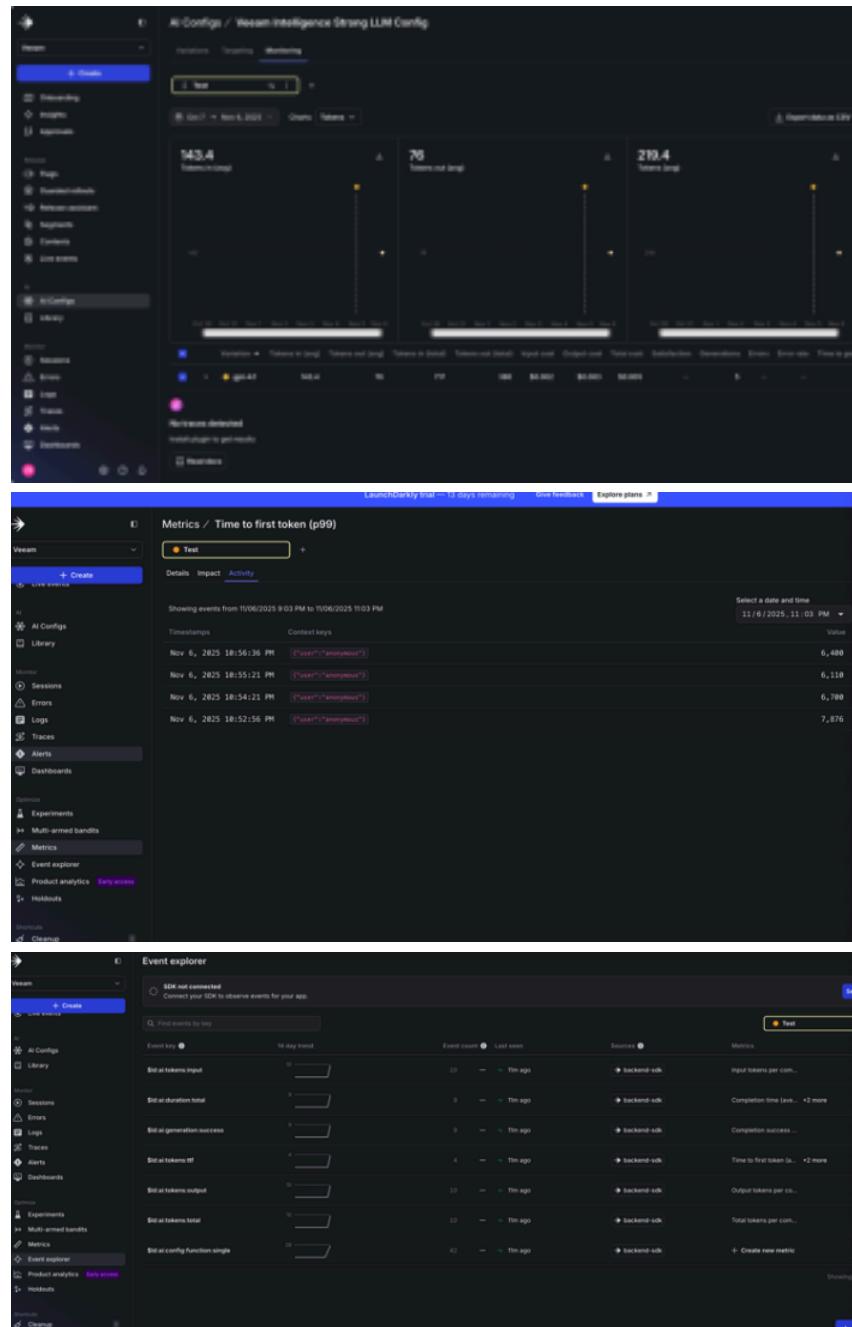
Input tokens - Tokens sent to model

Output tokens - Tokens generated by model

Duration - Time to generate

Time to first token - Latency metric

Errors - Failed generations



Experiments

I created an experiment to understand – can we quantify the difference in “time to first token” metric if we use a lighter weight

model like GPT 4.1-mini rather than GPT 4.1?

The screenshot shows the 'Experiments / Veeam Intelligence -- Model Latency' page. It displays a hypothesis: 'If I update my strong model from 4.1 to 4.1-mini, my latency decreases'. The experiment is set to 'Ready to start' with a status of 'Experiment design is complete'. The allocation split is 50% for 'gpt-4.1' and 50% for 'gpt-4.1-mini'. The interface includes sections for Hypothesis and metrics, Audience targeting, and Allocation.

Once the experiment was started, the targeting rules for my “strong” AI config were automatically updated according to the 50/50 allocation split for my experiment.

The screenshot shows the 'AI Configs / Veeam Intelligence Strong LLM Config' page. It displays a targeting configuration for the 'Test' environment. The traffic allocation is 50% for 'disabled', 50% for 'gpt-4.1', and 50% for 'gpt-4.1-mini'. A note states: 'If this rule is evaluated for context kinds other than [request] contexts, those contexts will be served.' The interface includes sections for Targeting configuration, AI Config, and Monitoring.

Note – when an AI config is being used as part of an experiment, you can't delete variations or change their value. This is smart, don't want to subvert the results

The screenshot shows the 'AI Configs / Veeam Intelligence Strong LLM Config' page, specifically the 'Variations' tab. A warning message states: 'Changes to variations may affect a running experiment. A running experiment is using this AI Config, so you cannot delete variations or change their values.' The interface includes sections for Variations, Targeting, Monitoring, and a 'Review and save' button.

At first, I noticed that the experiment wasn't tracking the results for requests and the AI config variations weren't being distributed 50/50. Rather, all 100% of requests were getting the gpt-4.1 model config. The issue was nuanced – my code wasn't using the correct “context kind”

The screenshot shows the Veeam AI Config interface. A targeting configuration for a 'Test' variation is displayed. The 'Default rule' is defined under the 'Experiment rule: Veeam Intelligence -- Model Latency'. The 'TRAFFIC ALLOCATION' section shows that 100% of [request] contexts will be in the experiment, with 0% disabled, 50% assigned to 'got-4.1', and 50% assigned to 'got-4.1-mini'. A note indicates that if the rule is evaluated for context kinds other than [request], those contexts will be served.

I updated my code to use the `request` context kind

The terminal window shows code changes in a file named `backend/product/plugins/vdc/backend/tar/launchdarkly_ai/config.py`. The changes involve updating the context kind from 'user' to 'request' and modifying session configurations. The commit message is:

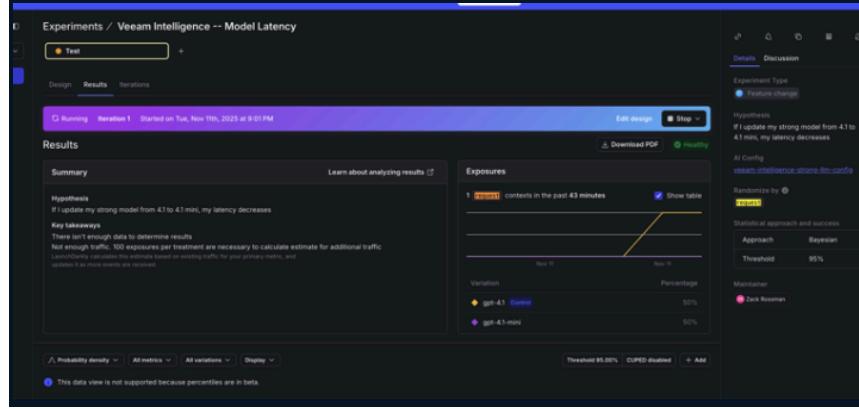
```
git commit -t "LaunchDarkly: Fix context kind for per-request experiments"
[...]
From the launch darkly doc: "If this rule is evaluated for context kinds other than request contexts, those contexts will be served." got-4.1

Is it possible my code isn't using "request" contexts, but a different context kind?
Let me check your context creation code.

Yes, that's exactly the issue! Looking at your LaunchDarkly configuration, the experiment is set up to split traffic for "request" context kind, but your code is creating a different context kind.
```

Hmmm... but even with the context kind set to `request`, Launch Darkly was still serving the default `gpt-4.1`

The screenshot shows the LaunchDarkly Contexts interface for the 'anonymous' context kind. It displays the 'Expected variations' section for AI Config, which lists two variations: 'gpt-4.1-mini' and 'gpt-4.1', both evaluated under the 'DEFAULT RULE'.



Debugging the issue with Warp AI agent

I see the issue! The context key is always "anonymous" (or the same user_id), which means Launch Darkly sees it as the same context across all requests. For per-request randomization in experiments, you need a unique key per request.

The experiment's traffic allocation uses the context key to determine which variation to serve. If the key is always the same, it will always hash to the same variation (gpt-4.1 in your case).

Here's the fix - generate a unique key per request:
Perfect! Now each request will have a unique UUID as the context key, which allows Launch Darkly's traffic allocation algorithm to properly randomize between the experiment variations.

What changed:

1. Unique key per request: Now using `uuid.uuid4()` to generate a unique key for each request
2. `user_id` as attribute: Moved `user_id` from the key to an attribute, so you can still target by user if needed in the future
3. Proper randomization: With unique keys, Launch Darkly can properly hash and distribute traffic according to your 50/50 experiment split

Why this fixes it:

Launch Darkly's traffic allocation uses the context key

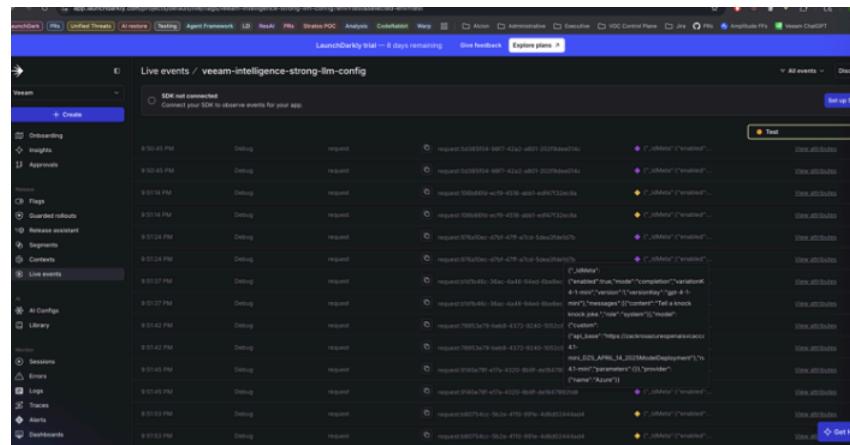
to hash and determine which variation to serve

When you were using the same key ("anonymous") for all requests, it always hashed to the same bucket → always gpt-4.1

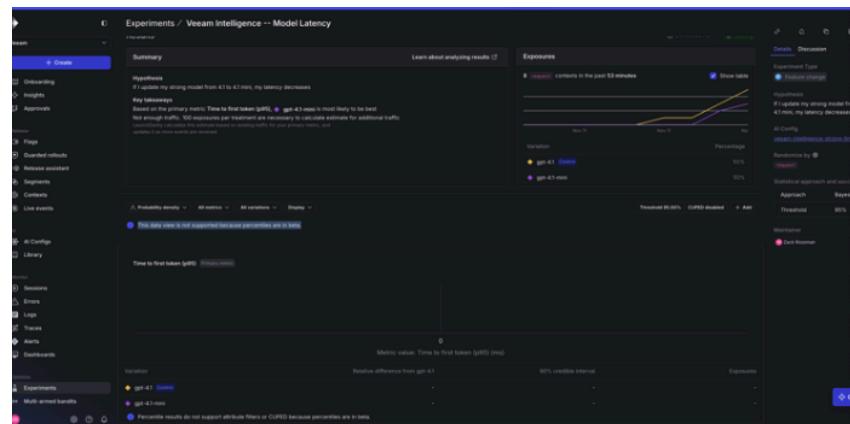
With unique keys per request, each request gets independently randomized → true 50/50 split

Now your experiment should properly distribute traffic between gpt-4.1 and gpt-4.1-mini!

The issue was I needed to add a unique **requestID** for each AI config evaluation. This fixes it and, indeed, now we can see an even split between requests 



Unfortunately, it seems the result view and the CUPED adjustment isn't supported for percentile analysis like p95 



Create metrics

Metrics measure audience behaviors affected by your experiments. You can use metrics to track all kinds of things, from how often end users access a URL to how long that URL takes to load a page. You can reuse metrics in multiple experiments, or create new ones for your feature change experiment.

To learn how to create your own new metric, read [Metrics](#). LaunchDarkly also automatically creates metrics for AI Configs. To learn more, read [Metrics generated from AI SDK events](#).

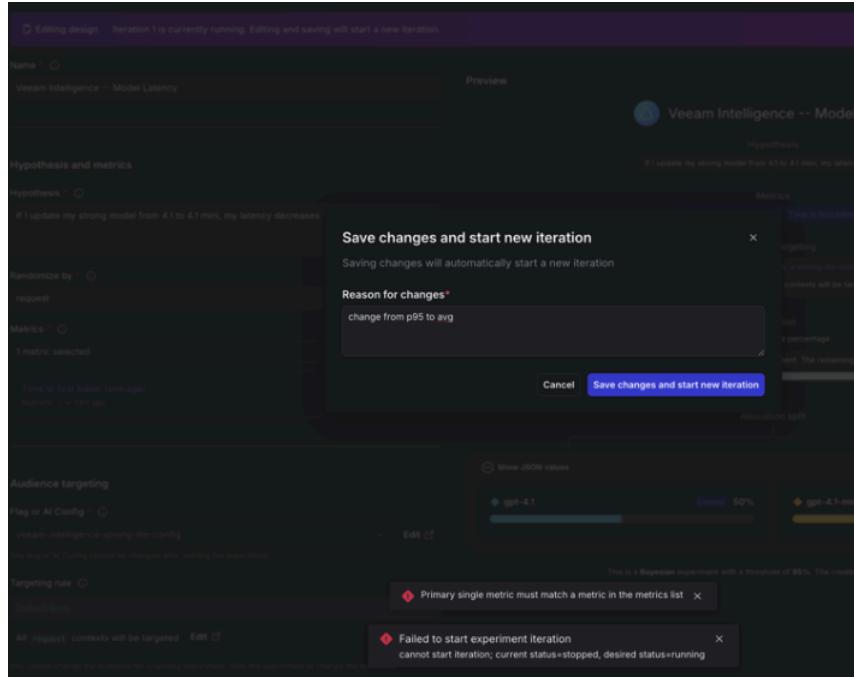
⚠ Percentile analysis methods for Experimentation are in beta

The default metric analysis method is "Average." The use of percentile analysis methods with LaunchDarkly experiments is in beta. If you use a metric with a percentile analysis method in an experiment with a large audience, the experiment results tab may take longer to load, or the results tab may time out and display an error message. Percentile analysis methods are also not compatible with [CUPED adjustments](#).

If you want to learn which variation performs better, it must be possible for that metric to be primary. [Creating feature change experiments](#) | [LaunchDarkly Documentation](#)

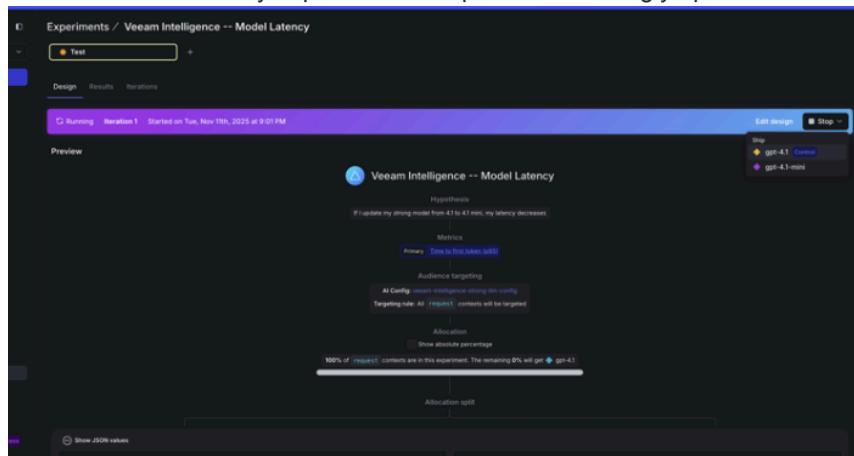
So I recreated the experiment where we're measuring **average** time to first token (rather than p95 time to first token) ⚠

I tried to create a new “iteration” of the existing experiment, but I had to start a new experiment altogether because the “primary single metric must match a metric in the metrics list”. So the issue was that I created the initial experiment with a percentage-based metric as the primary metric.



When you stop an experiment, Launch Darkly asks to “ship” a single variation. Is this a “feature” or a “bug”? I like it because it forces us to be decisive.

This is one example of how the Launch Darkly experimentation platform is strongly opinionated.



I recreated the experiment with time to first token (average) as the primary metric

They have a nice “health check” feature in on the right hand side of the experiment UI. This is a really thoughtful feature – helps us build confidence in our experiment setup and operations ☐

After 115 requests, 59 requests were handled by the gpt-4.1 variation and 54 requests were handled by the gpt-4.1- mini variation. We can see there's a 62% chance that gpt-4.1-mini is the superior variation if we want to pick the variation that reduces the average time to first token ☐

Evaluations

Online evals use AI judges, which are a type of AI Config that evaluates another AI Config's output in real time. Judges apply a consistent evaluation prompt and scoring framework to measure each response on key metrics such as accuracy, relevance, and toxicity.

Scores appear automatically on the Monitoring tab for each variation, alongside latency, cost, and satisfaction metrics. This provides a continuous signal of model performance with real users and data, enabling teams to detect regressions, improve reliability, and apply guardrails within the AI Config workflow.

Online evals differ from offline or pre-deployment testing. Instead of running evaluations manually in a sandbox or against datasets, they measure AI Config quality continuously in production.

Online evals work alongside observability. Observability helps you view model responses and routing details, while online evals provide quality scores you can use to trigger alerts, manage guarded rollouts, or run experiments.

Use online evals to:

Continuously monitor AI quality in production

Detect regressions immediately after a rollout

Automate rollbacks or alerts based on evaluation metrics

Compare prompt or model variations using live performance data

I attached all three judges to my AI configs

Judge	Event	Sampling percentage
AI Judge - Accuracy	Std ai judge accuracy	90%
AI Judge - Relevance	Std ai judge relevance	90%
AI Judge - Toxicity	Std ai judge toxicity	90%

The Launch Darkly judge evaluation capability is only supported in the [JS SDK](#). This is a bummer because we primarily [use Python and Go for our backend AI services](#). But this is a temporary limitation, estimated that the Python SDK will have support within 3 weeks (by [Dec 8, 2025](#)).

Observability

[LLM observability | LaunchDarkly | Documentation](#)

□ TODO

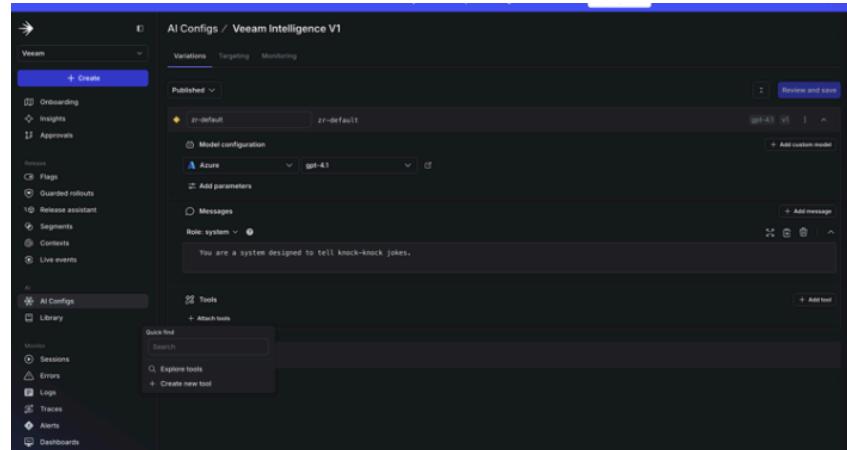
Tracing

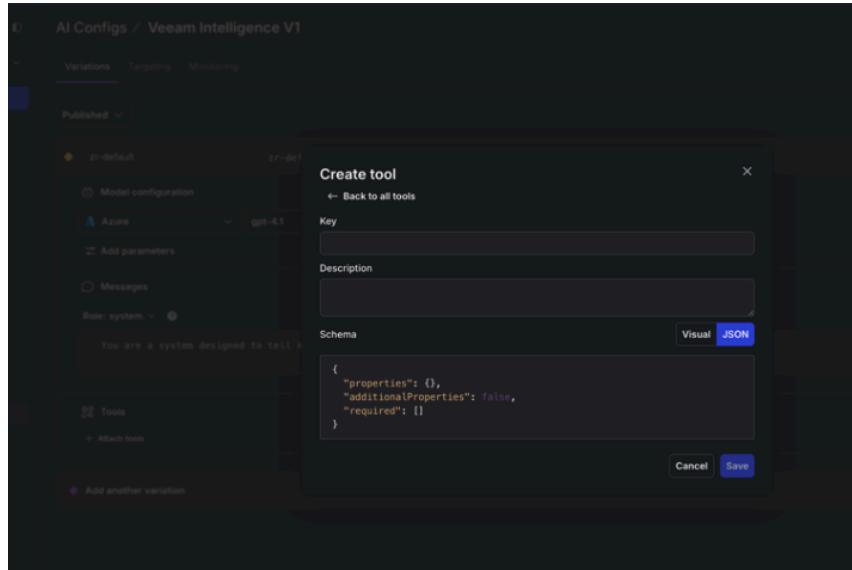
□ TODO

Creating Tools

[Tools in AI Configs | LaunchDarkly | Documentation](#)

IMO this feel like they're trying to do too much. Trying to apply "dynamic configuration" to every aspect of AI workflows, including agents and tools. Not sure if they're well-positioned in the stack to dictate agentic workflows and tool orchestration, probably better suited for dedicated AI agent framework code/platform. But I like that Launch Darkly has a pulse on AI/agentic trends.





TODO

Test dynamic prompts with placeholder values in the AI config prompt

POC Code

Here's the PR with the updates needed to integrate AI configs into Veeam Intelligence for VDC □

<https://github.com/veeam-ai/veeam-intelligence/pull/742> RESTRICTED CONTENT

Analysis

TODO Zack